

The Effect of Wealth on Individual and Household Labor Supply: Evidence from Swedish Lotteries[†]

By DAVID CESARINI, ERIK LINDQVIST, MATTHEW J. NOTOWIDIGDO,
AND ROBERT ÖSTLING*

We study the effect of wealth on labor supply using the randomized assignment of monetary prizes in a large sample of Swedish lottery players. Winning a lottery prize modestly reduces earnings, with the reduction being immediate, persistent, and quite similar by age, education, and sex. A calibrated dynamic model implies lifetime marginal propensities to earn out of unearned income from -0.17 at age 20 to -0.04 at age 60, and labor supply elasticities in the lower range of previously reported estimates. The earnings response is stronger for winners than their spouses, which is inconsistent with unitary household labor supply models. (JEL D14, J22, J31)

Understanding how labor supply responds to changes in wealth is critical when evaluating many economic policies, such as changes to retirement systems, property taxes, and lump-sum components of welfare payments. Because the income effect provides the link between uncompensated and compensated wage elasticities via the Slutsky equation, accurate estimates of how labor supply responds to wealth shocks are also valuable for obtaining credible estimates of compensated wage elasticities, which, in turn, are critical inputs in the theory of optimal taxation (Mirrlees 1971; Saez 2001) and studies of business cycle fluctuations (Prescott 1986; Rebelo 2005).

*Cesarini: Center for Experimental Social Science and Department of Economics, New York University, 19 W. 4th Street, 6FL, New York, NY 10012, NBER, and Research Institute of Industrial Economics (IFN) (email: david.cesarini@nyu.edu); Lindqvist: Department of Economics, Stockholm School of Economics, Box 6501, SE-113 83 Stockholm, Sweden, and IFN (email: erik.lindqvist@hhs.se); Notowidigdo: Northwestern University, 2001 Sheridan Road, Evanston, IL 60208, and NBER (email: noto@northwestern.edu); Östling: Institute for International Economic Studies, Stockholm University, SE-106 91 Stockholm, Sweden (email: robert.ostling@iies.su.se). This paper was accepted to the *AER* under the guidance of Hilary Hoynes, Coeditor. We thank André Chiappori, David Domeij, Trevor Gallen, John Eric Humphries, Erik Hurst, Edwin Leuven, Che-Yuan Liang, Jonna Olsson, Jesse Shapiro, John Shea, Johanna Wallenius, and seminar audiences at the AEA Annual Meeting, Bocconi, Chicago Booth, Cornell University, LSE, CREI-UPF, NBER Summer Institute, Northwestern University, the Rady School of Management, Stockholm University, and Uppsala University for helpful comments. We also thank Richard Foltyn, Victoria Gregory, My Hedlin, Renjie Jiang, Krisztian Kovacs, Odd Lyssarides, Svante Midander, and Erik Tengbjörk for excellent research assistance. This paper is part of a project hosted by IFN. We are grateful to IFN Director Magnus Henrekson for his strong commitment to the project and to Marta Benkestock for superb administrative assistance. The project is financially supported by three large grants from the Swedish Research Council (B0213903), Handelsbanken's Research Foundations (P2011:0032:1), and the Bank of Sweden Tercentenary Foundation (P15-0615:1). We also gratefully acknowledge financial support from the NBER Household Finance working group (22-2382-13-1-33-003), the NSF (1326635), and the Swedish Council for Working Life and Social Research (2011-1437). The authors declare that they have no relevant or material financial interests that relate to the research described in this paper.

[†]Go to <https://doi.org/10.1257/aer.20151589> to visit the article page for additional materials and author disclosure statement(s).

Despite a large empirical literature, consensus on the magnitude of the effect of wealth on labor supply is limited (Pencavel 1986; Blundell and MaCurdy 2000; Keane 2011; Saez, Slemrod, and Giertz 2012). Although some agreement exists among labor economists that large, permanent changes in real wages induce relatively modest differences in labor supply, Kimball and Shapiro (2008, p. 1) write that “there is much less agreement about whether the income and substitution effects are both large or both small.” The lack of consensus stems in part from the substantial practical challenges associated with isolating plausibly exogenous variation in unearned income or wealth, which is necessary to produce credible wealth-effect estimates. In this paper, we confront these challenges by exploiting the randomized assignment of lottery prizes to estimate the causal impact of wealth on individual- and household-level labor supply. Our work is most closely related to Imbens, Rubin, and Sacerdote’s (2001) survey of Massachusetts lottery players. Comparing winners of large and small prizes who gave consent to release their post-lottery earnings data from tax records, they estimate that around 11 percent of an exogenous increase in unearned income is spent on reducing pretax annual labor earnings.

Our lottery data are the same as in Cesarini et al. (2016) and have several advantages compared to previous lottery studies of labor supply (Kaplan 1985; Imbens, Rubin, and Sacerdote 2001; Furåker and Hedenus 2009; Larsson 2011; Picchio, Suetens, and van Ours 2017).¹ First, because we can effectively control for the number of lottery tickets bought, we only use the variation in lottery wealth that is truly exogenous. Second, the large size of the prize pool (approximately US\$650 million) allows us to precisely estimate heterogeneous effects across several subsamples. Because prizes vary in magnitude, we are also able to test for nonlinear effects.² Third, the lottery data is matched to Swedish high-quality administrative data that allow us to study labor market outcomes many years after winning the lottery essentially without attrition. Finally, we are able to address several frequently voiced concerns about the external validity of lottery studies.

In our reduced-form analyses of individual-level labor supply, we find winning a lottery prize immediately reduces earnings, with effects roughly constant over time and lasting more than ten years. Pretax earnings fall by approximately 1.1 percent of the prize amount per year. A windfall gain of 1 million SEK (about US\$140,000) thus reduces annual earnings by about 11,000 SEK, corresponding to 5.5 percent of the sample average.³ Adjustments of the number of hours worked account for the majority of the overall earnings response. Evidence of heterogeneous or nonlinear effects is scant, and winners are not more likely to change employers, industries, or occupations. We also find winning a lottery prize reduces self-employment income, which contrasts with several studies that find positive wealth shocks increase transition into self-employment (Holtz-Eakin, Joulfaian, and Rosen 1994; Lindh and Ohlsson 1996; Taylor 2001; Andersen and Nielsen 2012).

¹ Our work is also related to previous research that uses natural experiments such as policy changes or bequests to estimate the causal effect of wealth on labor supply (Bodkin 1959; Krueger and Pischke 1992; Holtz-Eakin, Joulfaian, and Rosen 1993; Joulfaian and Wilhelm 1994).

² The estimated effects in Imbens, Rubin, and Sacerdote (2001) are highly nonlinear and also somewhat sensitive to the small number of individuals in the sample who won prizes exceeding US\$2 million, as well as specifications that account for nonrandom survey nonresponse (Hirano and Imbens 2004).

³ All dollar amounts are converted using the January 2010 exchange rate (7.153 SEK per US\$).

We next estimate a simple dynamic labor supply model with a binding retirement age in order to extend the results beyond the first ten years following the prize event to estimate lifetime marginal propensities to earn (MPE) out of lottery wealth. The estimated model quantitatively accounts for our main reduced-form results. We account for the role of taxes by matching the model to the after-tax earnings response. The best-fit parameters imply the lifetime MPE varies with age and is strongest in the youngest winners, where our estimates suggest a lifetime MPE in the range -0.15 to -0.17 . Relying on the structural assumptions of the model, we also estimate key labor supply elasticities. The average uncompensated labor supply elasticity is close to zero, the individual-level compensated (Hicksian) elasticity is 0.10 , and the intertemporal (Frisch) elasticity is 0.14 . These estimates are in the lower range of previously reported estimates (Keane 2011; Reichling and Whalen 2012).

In our household-level analyses, we find that taking into account the labor supply of nonwinning spouses increases the estimated labor supply response by 23 percent. Our estimates are precise enough to reject both a zero effect on nonwinning spouses' earnings and the null hypothesis that the earnings responses of winning and nonwinning spouses are identical; we systematically find the winning spouse reacts more strongly. The latter result is inconsistent with unitary household labor supply models, which have the strong prediction that the observed labor supply responses to household wealth shocks should not depend on the identity of the lottery winner (Lundberg, Pollak, and Wales 1997).

Our finding that winners adjust labor supply more strongly than spouses complements a large empirical literature that uses labor supply data to test the exogenous income-pooling restriction of unitary models of the household (see the review by Donni and Chiappori 2011). As described in Lundberg and Pollak (1996, p. 145), an "ideal test of the pooling hypothesis would be based on an experiment in which some husbands and some wives were randomly selected to receive an income transfer." Our test comes close to these ideal conditions, and to our knowledge, we are the first to use random shocks to wealth from lottery prizes to directly test whether income is pooled when households make labor supply decisions.

The remainder of this paper is structured as follows. Section I describes the lottery data and reports the results from a randomization test. We discuss our empirical framework in Section II, and describe our measures of labor supply and report the individual-level empirical results in Section III. Section IV describes a dynamic life-cycle model and uses this model to estimate key labor supply elasticities. Section V reports household-level results and discusses how they inform household labor supply models. Section VI concludes the paper. We refer to our online Appendix for robustness tests and details regarding the data.

I. Lottery Samples

We construct our estimation sample by matching three samples of lottery players and their spouses to population-wide registers on labor market outcomes and demographic characteristics. The main challenge in lottery studies is that the number of lottery tickets held is correlated with the amount won. We overcome this challenge by using detailed institutional information about the lotteries to construct

TABLE 1—IDENTIFICATION ACROSS LOTTERY SAMPLES

	Time period (1)	Treatment variable (2)	Cells/fixed effects (3)	Number of cells (4)
PLS fixed-prize lottery	1986–2003	Sum of prizes	Prize draw \times number of prizes	206
PLS odds-prize lottery	1986–1994	Prize	Prize draw \times balance	1,620
Kombi lottery	1998–2010	Prize	Prize draw \times number of tickets \times age \times sex	260
Triss-Lumpsum	1994–2010	Prize	Year \times prize plan	18
Triss-Monthly	1997–2010	NPV of prize	Year \times prize plan	19

Notes: This table shows how we construct cells within which prizes are randomly assigned for the different lotteries. PLS odds prizes are only included for winners that win a single prize in a draw, and odds-prize cells with prizes totalling less than 100,000 SEK are excluded. NPV is net present value assuming an annual discount rate of 2 percent.

“cells” within which the amount won is random. We subsequently control for cell fixed effects in all analyses and thereby only use variation within cells to identify the causal effect of wealth. The cell construction is almost identical to Cesarini et al. (2016), but we reproduce the most essential details below for completeness.⁴ We describe the lotteries separately because the cells are constructed differently for each lottery. Table 1 summarizes the cell construction.

A. Prize-Linked Savings Accounts

The first sample is a panel of Swedish individuals who held prize-linked savings (PLS) accounts between 1986 and 2003. PLS accounts include a lottery element by randomly awarding prizes to some accounts rather than paying interest (Kearney et al. 2011). PLS accounts were initially subsidized by the Swedish government, but when the subsidies ceased in 1985, the government authorized banks to continue to offer PLS products. Two systems were introduced, one operated by the savings banks and one by the major commercial banks and the state bank. In the late 1980s, there were more than four million accounts in total, implying that about half of the Swedish population held a PLS account. We combine two sources of information from the PLS program run by the commercial banks, *Vinnarkontot* (“The Winner Account”). Our first source is a set of prize lists with information about all prizes won in the draws between 1986 and 2003. The prize lists were entered manually and contain information about prize amount, prize type (described below), and the winning account number, but not the identity of the winner. The second source is a large number of microfiche images with information about the account number, the account owner’s personal identification number (PIN), and the account balance of all eligible accounts participating in the draws between December 1986 and December

⁴The only difference compared to Cesarini et al. (2016) is that we use exact matching on age and sex for the Kombi lottery. A more detailed account of the institutional features of our three lottery samples, the processing of our primary sources of lottery data, data quality, and how cells were constructed is provided in the online Appendix to Cesarini et al. (2016).

1994 (the “fiche period”). By matching the prize-list data with the microfiche data, we are able to identify PLS winners between 1986 and 2003 who held an account during the fiche period.

In each draw, held every month throughout most of the studied time period, account holders were assigned one lottery ticket per 100 SEK in the account balance. Each lottery ticket had the same chance of winning a prize, so a higher account balance increased the chance of winning. There were two types of prizes: fixed prizes and odds prizes. Fixed prizes were regular lottery prizes that varied between 1,000 and 2 million SEK. The size of the prize did not depend on the account balance. In contrast, odds prizes paid either 1, 10, or 100 times the account balance (with the prize capped at 1 million SEK during most of the sample period).

We use different methods to construct the cells for the two types of prizes. For fixed-prize winners, we exploit the fact that the total prize amount is independent of the account balance among players who won the same number of fixed prizes in a particular draw. All winners who won an identical number of prizes in a draw are therefore assigned to the same cell. For example, one cell consists of 1,509 winners who won exactly one fixed prize in the draw of December 1990. We hence exclude account holders that never won from the sample, but we include fixed prizes won both during (1986–1994) and after (1995–2003) the fiche period. This identification strategy has been used in several previous studies of lottery players (Imbens, Rubin, and Sacerdote 2001; Hankins and Hoekstra 2011; Hankins, Hoekstra, and Skiba 2011).

For odds-prize winners, it is insufficient to match on the number of prizes won because the size of the prize is determined by the account balance. We therefore construct the cells by matching individuals who won exactly one odds prize to individuals with a similar account balance who also won one prize (odds or fixed) in the same draw. For example, one cell consists of an odds prize winner who won one odds prize in December 1990 and 19 winners of exactly one fixed prize in the same draw, all with account balances between 3,000 and 3,200 SEK. This matching procedure ensures that within a cell, the prize amount is independent of potential outcomes. Whenever a fixed-prize winner is matched to an odds-prize winner, the fixed-prize winner is excluded from the original fixed-prize cell. Thus, in the previous example from December 1990, none of the 19 fixed-prize winners in the odds-prize cell are included in the fixed-prize cell with 1,509 fixed-prize winners. An individual is hence assigned to at most one cell in a given draw, but because players can win in several draws, they will also be included in cells corresponding to other draws in which they won. We only include odds prizes won during the fiche period (1986–1994) because we do not observe account balances after 1994. To keep the number of cells manageable, we exclude all odds-prize cells in which the total amount won is below 100,000 SEK.

B. The Kombi Lottery

Our second sample consists of about half a million individuals who participated in a monthly ticket-subscription lottery called *Kombilotteriet* (“Kombi”). The proceeds from Kombi go to the Swedish Social Democratic Party, Sweden’s main political party during the postwar era. Subscribers choose their desired number of subscription tickets and are billed monthly. Our dataset contains information about

all draws conducted between 1998 and 2010. For each subscriber, the data contain information about the number of tickets held in each draw and information about prizes exceeding 1 million SEK. The Kombi rules are simple: two individuals who purchased the same number of tickets in a given draw face the same probability of winning a large prize. To construct the cells, we match each winning player to (up to) 100 randomly chosen nonwinning players of the same age and sex with an identical number of tickets in the month of the draw. Complete random assignment of prizes within cells requires that controls are drawn with replacement from the set of potential controls. Players who win in one draw may therefore be used as controls in a different draw, and some individuals are used as controls in multiple draws. We exclude four winners who could not be exactly matched to any controls.

C. Triss Lotteries

Triss is a scratch-ticket lottery run by the Swedish government-owned gaming operator Svenska Spel since 1986. Triss lottery tickets are sold in most convenience and grocery stores throughout the country. The sample we have access to consist of two categories of winners: Triss-Lumpsum and Triss-Monthly. Winners of either type of prize are invited to participate in a morning TV show. At the TV show, Triss-Lumpsum winners draw a new scratch-off ticket from a stack of tickets with prizes varying from 50,000 to 5 million SEK. Triss-Monthly winners draw two lottery tickets independently. One ticket determines the size of a monthly installment and the second its duration. The prize duration varies from 10 to 50 years, and the monthly installments from 10,000 to 50,000 SEK. We convert the installments in Triss-Monthly to their present value using a 2 percent annual discount rate.⁵ Our data includes participants in Triss-Lumpsum and Triss-Monthly prize draws between 1994 and 2010 (the Triss-Monthly prize was not introduced until 1997).⁶ We exclude about 10 percent of the lottery prizes for which the data indicate the player shared ownership of the ticket.

The prize plans used in the TV show are subject to occasional revision, but for a given prize plan and conditional on qualifying for the TV show, the amount won in Triss-Lumpsum or Triss-Monthly is random. We therefore assign players to the same cell if they won one prize (of a given type) in the same year and under the same prize plan. No suitable controls exist for players that won more than one prize of the same type within a year and under the same prize plan, and a few such cases have been excluded.

D. Estimation Sample

Merging the three lotteries gives us a sample of 435,966 observations. Primarily because many PLS lottery players win small prizes several times, these observations correspond to 334,532 unique individuals. To arrive at our estimation sample, we

⁵ We set the discount rate to match the real interest rate in Sweden, which, according to Lagerwall (2008), was 1.9 percent during 1958–2008.

⁶ The original data file did not include information on personal identification numbers (PINs), but using information on name, age, and address, 98.7 percent of the lottery show participants could be reliably identified.

TABLE 2—DISTRIBUTION OF PRIZES

SEK	Pooled sample		Individual lottery samples							
			PLS		Kombi		Triss-Lumpsum		Triss-Monthly	
	Count (1)	Share (2)	Count (3)	Share (4)	Count (5)	Share (6)	Count (7)	Share (8)	Count (9)	Share (10)
0 to 1K	23,910	0.10	0	0.00	23,910	0.99	0	0.00	0	0.00
1K to 10K	201,600	0.82	201,600	0.92	0	0.00	0	0.00	0	0.00
10K to 100K	16,284	0.07	15,376	0.07	0	0.00	908	0.28	0	0.00
100K to 500K	3,656	0.01	1,632	0.01	0	0.00	2,024	0.62	0	0.00
500K to 1M	355	0.00	195	0.00	0	0.00	160	0.05	0	0.00
≥1M	1,470	0.01	471	0.00	262	0.01	168	0.05	569	1.00
Total	247,275		219,274		24,172		3,260		569	

Notes: This table reports the distribution of lottery prizes for the pooled sample and the four lottery subsamples. The development of the prize distribution over time for each lottery is shown in online Appendix Figure A2.

first exclude individuals who (i) died the same year they won; (ii) lack information on basic socioeconomic characteristics in public records; or (iii) have no recorded income in any year up to 10 years after winning, leaving us with a sample of 426,598 observations. We further restrict attention to players who were between age 21 and 64 at the time of the win, which reduces the sample to 249,402 observations. Finally, we drop cells without variation in the amount won and end up with an estimation sample of 247,275 observations (200,937 individuals).

E. Prize Distribution

Table 2 shows the distribution of prizes in the pooled sample and for each lottery separately. All lottery prizes are net of taxes and expressed in units of year-2010 SEK. Among the 247,275 lottery players in our sample, less than 9 percent won more than 10,000 SEK (US\$1,400). Yet in total more than 5,500 prizes are in excess of 100,000 SEK (US\$14,000) and almost 1,500 prizes are in excess of 1 million SEK (US\$140,000). To put these numbers into perspective, the median disposable income in a representative sample of Swedes in 2000 was 170,000 SEK. The total prize amount in our pooled sample is 4,662 million SEK (about US\$650 million). PLS and Triss-Monthly each account for 36 percent of the total prize amount, Triss-Lumpsum account for 21 percent, and Kombi 7 percent.⁷ The fact that most prizes are small does not imply our estimates are mostly informative about the marginal effects of wealth at low levels of wealth. The reason is that most of the identifying variation in our data comes from within-cell comparisons of winners of small or moderate amounts to large-prize winners.

F. Internal and External Validity

Internal Validity.—Key to our identification strategy is that the variation in amount won within cells is random. If the identifying assumptions underlying the

⁷ Online Appendix Figure A2 shows the total prize amount is quite stable over time for all lotteries except PLS, where the prize sum falls from the late 1980s onward. Figure A2 also shows that, for each lottery, the proportion of the prize sum awarded within a certain prize range is quite stable over time.

lottery cell construction are correct, then characteristics determined before the lottery should not predict the amount won once we condition on cell fixed effects, because, intuitively, all identifying variation comes from within-cell comparisons. To test for violation of conditional random assignment, we therefore run the regression

$$(1) \quad L_{i,0} = \mathbf{X}_{i,0}\boldsymbol{\eta} + \mathbf{Z}_{i,-1}\boldsymbol{\theta} + \varepsilon_{i,0},$$

where $L_{i,0}$ is the total amount won, $\mathbf{X}_{i,0}$ is a vector of cell fixed effects, and $\mathbf{Z}_{i,-1}$ is a vector of baseline controls. In our tests for random assignment, the controls are indicator variables for sex, born in the Nordic countries, college completion, pretax annual labor earnings, and a third-order polynomial in age. All time-varying baseline controls are measured in the year prior to the lottery. We estimate this equation for the pooled sample and for each lottery separately. For the pooled sample, we also estimate the equation without cell fixed effects. The results are consistent with the null hypothesis that wealth is randomly assigned conditional on the fixed effects (see online Appendix Table A1).

External Validity.—An important concern about lottery studies is that lottery players may not be representative of the general population. We therefore compare the demographic characteristics of players in each of our lottery samples to random population samples drawn in 1990 and 2000. Men are over-represented in the Kombi sample (60.9 percent) and the average age in the lottery sample (48.6 years) is about 7 years older than in the population (see online Appendix Figure A1 for the full age and sex distribution). Consequently, characteristics that vary substantially between the sexes or over the life cycle (such as income) will differ between players and the population. To adjust for such compositional differences, we reweight the representative samples to match the age and sex distribution of the lottery winners. Compared to the reweighted representative samples, lottery players are more likely to be born in the Nordic countries and (except for the PLS lottery) have lower levels of education, but are quite similar with respect to income and marital status (see online Appendix Table A2 for detailed results). A related concern is that, even though lottery players may be similar to the population at large, lottery prizes constitute a specific type of wealth shock that cannot be generalized to other types of wealth. Although we cannot rule out this concern completely, online Appendix Figure A5 shows household wealth dissipates slowly with time since winning the lottery. Moreover, as shown below, the winners' estimated labor supply response fits the predictions of standard life-cycle models fairly well irrespective of the type of lottery (PLS, Kombi, or Triss) or mode of payment (Triss-Lumpsum or Triss-Monthly).

II. Estimation Strategy

Normalizing the time of the lottery to $t = 0$, our basic estimating equation is

$$(2) \quad y_{i,t} = \beta_t L_{i,0} + \mathbf{X}_{i,0}\boldsymbol{\delta}_t + \mathbf{Z}_{i,s}\boldsymbol{\gamma}_t + \varepsilon_{i,t},$$

where $y_{i,t}$ is individual i 's year-end outcome of interest measured at time $t = 0, 1, \dots, 10$, $L_{i,0}$ is the lottery prize won, and $\mathbf{X}_{i,0}$ is a vector of cell fixed

effects. The term $\mathbf{Z}_{i,s}$ includes the lagged outcome of the dependent variable plus the same vector of pre-lottery characteristics as in equation (1) measured s years prior to winning the lottery. The identifying assumption is that $L_{i,0}$ is independent of potential outcomes conditional on $\mathbf{X}_{i,0}$. We control for the pre-lottery characteristics in $\mathbf{Z}_{i,s}$ to improve precision. We estimate equation (2) by OLS and cluster standard errors at the level of the individual.

We report our main results in two formats. First, we often summarize our results by plotting the coefficient estimates $\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_{10}$ in a figure with time (in years) on the horizontal axis. These figures show the dynamic effects of a $t = 0$ wealth shock on the labor supply outcome of interest at $t = 0, 1, \dots, 10$. To verify the absence of differences in pre-treatment characteristics of players who won large or small prizes, the figures also include estimates for two or four pre-lottery years, which should not be significantly different from zero under our identifying assumption. In these regressions, the time-varying controls in $\mathbf{Z}_{i,s}$ are measured one year prior to the first estimate shown in the figure (i.e., s is -3 or -5 depending on the first estimate shown).

Second, we report estimates from a modified version of equation (2) in which we include all person-year observations available for $t = 1, \dots, 5$ and use base-line controls measured in the year prior to the lottery event (i.e., $s = -1$). We also impose the restriction that $\beta_t = \beta$, which, as we show below, is motivated in part by our empirical evidence that the response to wealth shocks is near-immediate and quite stable over time for most outcomes we consider. We refer to these estimates as five-year estimates. The five-year estimates allow us to improve precision and present our findings in a parsimonious way.

Because small average effects could mask large effects in certain subpopulations, we also test for heterogeneous effects. In these analyses, we interact the lottery prize, $L_{i,0}$, the vector of cell fixed effects, \mathbf{X}_i , and the controls, $\mathbf{Z}_{i,-s}$, with the subpopulation indicator variable of interest, thereby leveraging only within-cell variation to estimate treatment-effect heterogeneity.

III. Individual-Level Analyses

In this section, we analyze individual-level responses to lottery wealth shocks. We begin by describing and analyzing a number of annual earnings measures. Next, we decompose the total wealth effect on earnings into extensive- and intensive-margin adjustments, and into hours and wage changes. The section concludes with analyses in which we test for treatment-effect heterogeneity and nonlinear effects.

All analyses below are restricted to labor supply outcomes observed from 1991 until 2010 (the last year for which we have data).⁸ The annual income measures we use are based on population-wide registers that contain information originally collected by the tax authorities. All income variables are winsorized at the 0.5th and 99.5th percentile and are expressed in units of year-2010 SEK.

⁸ Sweden underwent a major tax reform in 1990–1991. Before 1991, capital and labor incomes were taxed jointly and taxes were strongly progressive, which complicates the analysis of wealth effects.




FIGURE 1. EFFECT OF WEALTH ON INDIVIDUAL EARNINGS

Notes: This figure reports estimates obtained from equation (2) estimated in the pooled lottery sample with pretax labor earnings as the dependent variable. A coefficient of 1.00 corresponds to an increase in annual earnings of 1 SEK for each 100 SEK won. Each year corresponds to a separate regression and the dashed lines show 95 percent confidence intervals.

A. *Effect of Wealth on Annual Earnings*

Our primary earnings measure is pretax labor earnings, a composite variable derived almost entirely from three sources of income: annual wage earnings, income from self-employment, and income support due to parental leave or sickness absence. Figure 1 depicts the estimated effect of wealth on our primary outcome for $t = -4, -3, \dots, 10$ along with 95 percent confidence intervals. Consistent with the identifying assumption of conditional random assignment of lottery prizes, the point estimates in the years prior to winning are statistically indistinguishable from zero. The effect of lottery wealth is near-immediate, modest in size, and quite stable over time.⁹ The tendency for the effect to decline over time vanishes if we restrict the sample to individuals who were below age 55 at the time of winning and who therefore had at least 10 years left to age 65, the modal retirement age in Sweden (see Figure 3, panel B).¹⁰ As discussed in Section IV, a stable response over

⁹For the average winner, labor earnings in $t = 0$ include income from six months prior to and six months after the lottery draw, so the fact that the point estimate at $t = 0$ is about half the $t = 1$ estimate suggests that lottery players adjust labor supply quickly after winning the lottery.

¹⁰Because we limit the sample to labor earnings measured in 1991–2010 and the sample consists of individuals who won the lottery in 1986–2010, the composition of the pooled sample in Figure 1 changes somewhat with t . For example, an individual who won the lottery in 1986 will not enter the data until $t = 5$. Conversely, an individual who won in, say, 2010, will exit the data at $t = 1$. In online Appendix Section 4, we show the time pattern of labor supply responses looks quite similar up until $t = 10$ when we hold the sample fixed. The data indicate larger responses after $t = 10$, but due to the smaller sample size, we rely on the model in Section IV instead of these estimates to make inferences about long-term effects of lottery wealth on labor supply.

TABLE 3—EFFECT OF WEALTH ON ANNUAL INCOME

	Pretax labor earnings $\approx (2) + (3) + (4)$ (1)	Wage earnings (2)	Self-employment income (3)	Income support (4)	Production value = (1) + SSC (5)
Effect (100 SEK)	-1.066	-0.964	-0.051	-0.016	-1.412
SE	(0.148)	(0.151)	(0.030)	(0.030)	(0.199)
<i>p</i>	[<0.001]	[<0.001]	[0.088]	[0.594]	[<0.001]
Mean (SEK)	194,505	175,608	6,598	7,327	258,757
Effect/mean	-5.48	-5.49	-7.72	-2.15	-5.46
Observations	244,826	244,826	244,826	244,826	244,826

	Taxable labor income $\approx (1) + (7) + (8)$ (6)	Unemployment benefits (7)	Pension income (8)	Taxes (excl. SSC) (9)	After-tax labor income = (6) - (9) (10)
Effect (100 SEK)	-0.890	0.024	0.114	-0.313	-0.576
SE	(0.131)	(0.024)	(0.070)	(0.054)	(0.080)
<i>p</i>	[<0.001]	[0.317]	[0.104]	[<0.001]	[<0.001]
Mean (SEK)	244,837	4,615	29,973	83,134	161,703
Effect/mean	-3.64	5.20	3.80	-3.76	-3.56
Observations	244,826	244,826	244,826	244,826	244,826

Notes: This table reports five-year estimates (per 100 SEK won) obtained by estimating equation (2) for various annual income measures as dependent variables. Pretax labor earnings (column 1) is approximately the sum of income from wages (column 2), self-employment (column 3), and parental leave and sickness payments (column 4). Taxable labor income is approximately the sum of pretax labor earnings (column 1), unemployment benefits (column 7), and pensions (column 8). For details on the calculations of taxes owed, see online Appendix Section 3.2. SSC: social security contributions paid by employer. Effect/mean: effect of 1 million SEK relative to the mean of the dependent variable (in percent). Standard errors are clustered at the level of the player.

time is consistent with a canonical life-cycle model where the discount factor equals the interest rate. The lottery variable is measured in units of 100 SEK, so the coefficient estimates of approximately -1 means winners reduce their annual earnings by ~ 1 percent of the prize amount per year. To help interpret the magnitude of these point estimates, a 1 million SEK prize (US\$140,000) reduces earnings by about 10,000 SEK, which corresponds to 5.5 percent of the annual average earnings in our sample.

A more detailed picture of the labor supply response is provided in Table 3, which shows the five-year estimates for pretax labor earnings along with a range of additional income measures. We begin in the upper panel, which reports results for the pretax labor earnings variable and the three variables from which it is derived: wage earnings, self-employment income, and income support due to parental leave and sick leave. The results are shown in columns 1 to 4. Unsurprisingly, nearly all of the overall effect on pretax labor earnings (-1.066) is accounted for by reductions in wage earnings (-0.964), with self-employment income also contributing modestly (-0.051) to the overall decline.¹¹ Yet the effect on self-employment relative

¹¹ The three coefficient estimates in columns 2 to 4 do not add up exactly to the coefficient estimate in column 1, because labor earnings include some other minor forms of income support not included in column 4. However, the correlation between labor earnings and the sum of wage earnings, self-employment income, and our measure of income support is 0.99.

to baseline is actually larger than for wage earnings: a 1 million SEK windfall gain reduces self-employment income by 7.7 percent of the annual average compared to 5.5 percent for wage earnings. The reduction in self-employment income is at odds with previous findings that windfall gains increase self-employment (Holtz-Eakin, Joulfaian, and Rosen 1994; Lindh and Ohlsson 1996; Taylor 2001; Andersen and Nielsen 2012). The effect on income support is very small (-0.016) and not statistically significant.

The pretax labor earnings measure includes income taxes, but not so-called social security contributions (SSC) paid by the employer. These contributions are partly taxes and partly benefits that accrue to the employee, for example, in the form of higher pension income in the future. Pretax labor earnings plus SSC represent the employers' total labor cost and can hence be interpreted as a measure of total production value. Column 5 of Table 3 shows the estimated impact of wealth on earnings plus SSC. According to our estimate, a 100 SEK windfall reduces the total production value by 1.41 SEK per year in the first five post-lottery years.

We also examine how lottery wealth affects after-tax income. In Sweden, labor market earnings are taxed jointly with unemployment benefits and pension income, so we use a measure of taxable labor income that includes all three sources of income. Column 6 shows the estimated impact on this measure (-0.890) is smaller than the impact on our primary earnings measure in column 1 (-1.066). The difference arises because lottery wealth causes a small increase in pension income (column 7) and unemployment benefits (column 8), and these benefits partly offset the reduction in labor earnings.¹² We use detailed information about the Swedish tax system to calculate the implied after-tax labor income for each winner. As shown in column 10, the estimated effect on after-tax income (-0.576) is substantially smaller than the effect on total production value in column 5 (-1.412).¹³ The difference reflects the wedge induced by Sweden's extensive tax and transfer system.

How large is the after-tax labor supply response from a life-cycle perspective? The average winner in our sample is 48.6 years old and thus has roughly 16.4 years of work left before the typical retirement age of 65. Ignoring discounting and assuming a constant effect of wealth on labor supply, lifetime after-tax income decreases by $0.576 \times 16.4 = 9.44$ SEK per 100 SEK won. This approximation is a simple estimate of the lifetime MPE out of unearned income. Relating the labor supply response to average total lifetime wealth before the win (wealth and future earnings and pensions) of approximately 4.7 million SEK allows us to get a rough estimate for the labor supply elasticity with respect to lifetime income.¹⁴ For such a winner, a 1M prize increases lifetime wealth by $1/4.7 = 21$ percent and decreases after-tax labor income by 3.6 percent, implying an elasticity of about -0.17 . This wealth elasticity is within the range of income elasticities reviewed by the Congressional Budget Office

¹²The estimate in column 6 is not exactly equal to the sum of the estimates in columns 1, 7, and 8 because other minor differences exist between pretax labor earnings and taxable labor income we have not taken into account here.

¹³Including the value of future benefits (notably pensions) implicit in SSC in our after-tax income measure increases the estimated effect to -0.624 .

¹⁴In addition to the assumptions made above, we assume wage and pension growth of 2 percent per annum, a post-tax income of 147,000 per year until retirement at age 65, a retirement replacement rate of 70 percent, a remaining life span of 30 years, and pre-win wealth of 0.9 million SEK.

TABLE 4—MARGINS OF ADJUSTMENT

	<i>Panel A. Extensive margin</i>			<i>Panel B. Retirement</i>		<i>Panel C. Hours and wages</i>	
	Labor earnings (1)	Wage earnings (2)	Self-empl. income (3)	Pension income (4)	Quit work before 65 (5)	Weekly hours (6)	Monthly wage (7)
Effect (million SEK)	−2.015	−2.241	−0.139	0.951	3.302	−1.282	−147.3
SE	(0.435)	(0.473)	(0.202)	(0.658)	(1.420)	(0.247)	(84.2)
<i>p</i>	[<0.001]	[<0.001]	[0.491]	[0.148]	[0.020]	[<0.001]	[0.080]
Mean	0.78	0.71	0.05	0.50	0.51	32.8	22,999
Effect/mean	−2.60	−3.15	−2.63	1.91	6.43	−3.91	−0.64
Observations	244,826	244,826	244,826	89,980	74,718	108,919	108,919

Notes: All columns except column 5 report five-year estimates (per 1 million SEK won) obtained by estimating equation (2) in the pooled lottery sample. In columns 1 to 4, each dependent variable is equal to 1 if annual income exceeded 25,000 SEK, and 0 otherwise. Column 5 reports the estimate obtained by estimating equation (2) with a dependent variable that takes the value 1 if annual labor earnings was below 25,000 SEK at both age 64 and 65, and 0 otherwise. The sample in panel B is restricted to individuals aged 55 or above at the time of the lottery event. The variables in panels A and B are scaled so that a coefficient of 1.00 implies a 1 percentage point increase in (non) participation per million SEK won. In panel C, the dependent variables are the imputed number of hours worked per week (column 6) and the pretax monthly wage (column 7). Effect/mean: Effect of 1 million SEK relative to the mean of the dependent variable (in percent). Standard errors are clustered at the level of the lottery player.

(CBO), which found estimates between -0.2 and 0 (CBO 1996; McClelland and Mok 2012).¹⁵

B. Margins of Adjustment

In this section, we decompose the overall effect on annual earnings into various margins of adjustment. We begin by estimating extensive-margin and retirement responses, and then turn to estimating the effect on wages and hours worked. To understand potential mechanisms, we also analyze whether lottery winners adjust their labor supply by changing occupations, employers, workplaces, industries, or location of work. The key results from our analyses are reported in Table 4 and Figure 2.

We first estimate the effect of wealth on several extensive-margin indicator variables generated from labor earnings, wage earnings, and self-employment income. For each category, we define an indicator equal to 1 if annual income exceeded 25,000 SEK (US\$3,500) in a given year, and 0 otherwise. The five-year estimates from these analyses are shown in panel A of Table 4. We scale the treatment variable so that a coefficient of 1.00 means 1M increases participation probability by 1 percentage point. We also report coefficient estimates normalized by the baseline participation probability.

Figure 2, panel A, shows winning the lottery reduces labor force participation by about 2 percentage points per 1 million SEK won up until five years after the

¹⁵Note the wealth elasticity we calculate is not the exact same concept as the income elasticity estimates reviewed in the CBO reports. In those reports, the income-elasticity estimates represent the elasticity of hours worked with respect to total after-tax income, holding constant the marginal after-tax wage rate. By contrast, our reduced-form estimate of the effect of lottery wealth includes both the effects on hours and wages, and so it does not hold the wage constant.

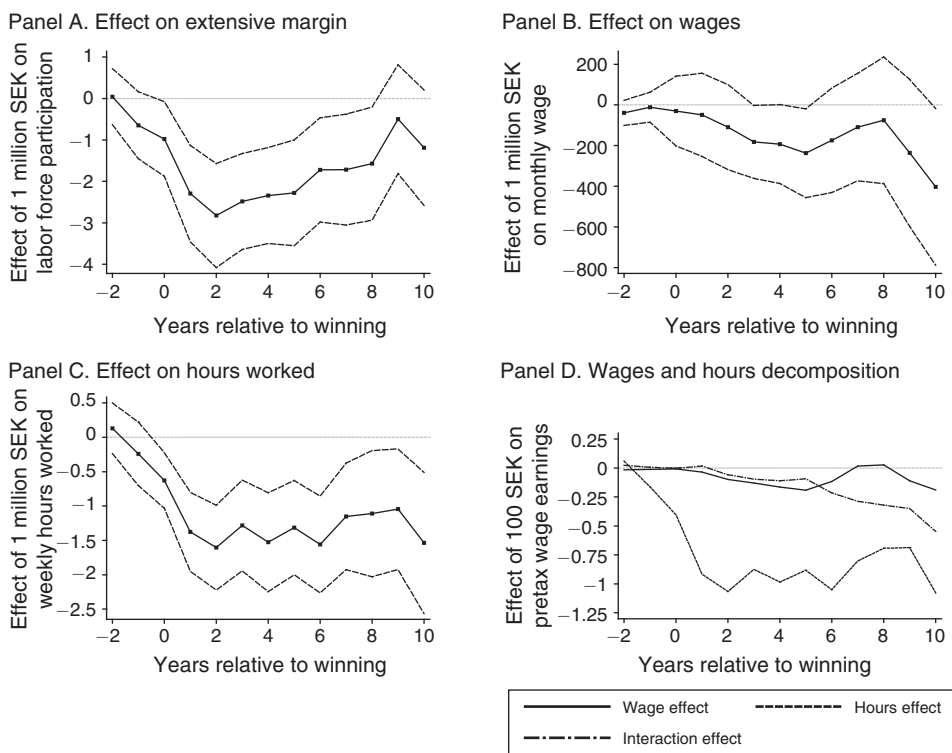


FIGURE 2. MARGINS OF ADJUSTMENT

Notes: Panels A–C report estimates obtained from equation (2) estimated for some of the different margins of adjustment discussed in Section IIIB. Each year corresponds to a separate regression. The dashed lines in panels A–C display 95 percent confidence intervals. Panel D reports the wage-hours decomposition described by equation (3) for each year separately, but using lagged values from $t = -3$ rather than $t = -1$.

win, after which the effect declines.¹⁶ The five-year estimates in panel A of Table 4 show the reduction in participation (–2.0 percentage points per 1 million SEK won) is almost entirely due to a fall in the probability of wage labor (–2.2 percentage points) rather than self-employment income (–0.1 percentage points). Yet because the baseline incidence of self-employment is lower, the responses are similar in relative terms (–3.1 percent and –2.6 percent).

The estimated effects for the extensive margin imply much of the labor supply response occurs on the intensive margin, in the form of lower wages or fewer hours. Under the assumption that average wage earnings of workers who leave the labor force due to winning the lottery equal the sample average, the extensive margin accounts for about 40 percent of the five-year labor supply response.¹⁷ Because the

¹⁶ Figure 2, panel A, might give the impression that lottery winners have different trends in labor force participation prior to winning. However, the $t = -1$ estimate is not significantly different from zero ($p = 0.116$), and since Figure 2 and Figure 3 report estimates for several outcomes and groups, it is not surprising if some estimates are nonzero. To minimize the influence of pre-win differences in outcomes, all results reported in tables control for the dependent variable measured at $t = -1$.

¹⁷ Because we do not observe the counterfactual earnings level of workers whose choice regarding whether to leave or enter the labor market was influenced by the lottery win, the decomposition into extensive- and intensive-margin responses is only suggestive. A more elaborate analysis where we estimate the effect of winning

estimated effect on the extensive margin declines faster than the overall labor supply response, the importance of intensive-margin adjustments increases with time from the lottery event.¹⁸

Column 4 of Table 4 reports the effect on receiving pension income above 25,000 SEK for winners aged at least 55 at the time of winning. We estimate a small positive, but statistically insignificant, effect of winning the lottery on the probability of receiving pension income. Because it is not possible to claim pension benefits early for many workers (see online Appendix Section 6 for details about the Swedish pension system), some may retire without claiming pension benefits. To investigate this possibility, we estimate the effect on quitting the labor force prior to age 65 (defined as earning less than 25,000 SEK at both age 64 and 65). We restrict the sample to winners aged at least 55 whom we can follow up until age 65. We find winning 1M increases the probability of early retirement defined accordingly by 3.3 percentage points.

Our third set of analyses focuses on wages and hours worked. We supplement the register-based variables with information from Statistics Sweden's annual wage survey. The survey asks employers to supply information about each employee's full-time equivalent monthly wage and the number of hours the individual is contracted to work.¹⁹ The survey has incomplete coverage of the private sector and covers 57 percent of the working population (those with wage earnings above 25,000 SEK) the year before the lottery win. The survey sample is not fully representative of the population of lottery winners.²⁰ In our main analyses, we impute information from adjacent years, increasing coverage to 67 percent of the working population.²¹ Even after imputation, the survey measure on contracted hours has two potential problems. First, modest adjustment of hours worked on a number of margins, such as sick leave, unpaid vacation, and over-time, may not induce changes in contracted hours. Second, because the survey only covers the employed, individuals who are induced by the lottery wealth shock to leave their job are absent from the survey, creating a potential selection problem. To mitigate these problems, we use the register-based data on wage earnings to calculate an earnings-based measure for weekly hours worked:

$$\text{Weekly hours} = 40 \times \frac{\text{Annual wage earnings}}{12 \times \text{Contracted monthly wage}}.$$

on entry and exit probabilities separately, and calculate counterfactual earnings based on pre-win entrants and exiters, shows the extensive margin accounts for about a third of the overall response in the first five years after the lottery win.

¹⁸ Applying the same back-of-the-envelope calculation as above, the share attributed to the extensive margin goes from around 40 percent in the first five years after the lottery to 24 percent ten years after the lottery.

¹⁹ The wage survey also contains a measure of actual hours worked during September–October every year, but this variable is only available from 1996 for a smaller sample, and suffers from the same selection problem as contracted hours.

²⁰ Lottery players in the wage-hours sample are about two years younger and have 19 percent higher earnings compared to the baseline sample. The effect of winning on labor earnings is similar in the first five years after the lottery. The five-year estimate is -1.064 for the wage-hours sample compared to -1.066 in the full sample, but the response in later years is somewhat larger in the wage-hours sample (see online Appendix Figure A12).

²¹ In our baseline specification, we impute observations for year t from up to $t - 3$ and $t + 3$ when data closer to t are unavailable. However, we never impute observations for post-win years from pre-win years, or vice versa. Further details about the imputation procedure are presented in online Appendix Section 5.

Because wage earnings are observed for the full sample in each year, the earnings-based measure will capture hours worked quite accurately also for workers who work few hours, as long as we are able to impute the wage from adjacent years.²²

The five-year estimates of the impact of wealth on earnings-based hours and monthly wages are shown in panel C of Table 4. Column 6 shows a 1 million SEK prize reduces (earnings-based) weekly hours by 1.3, corresponding to 4 percent of an average workweek. The estimate in column 7 shows the estimated impact on the pretax monthly wage (rescaled to its full-time equivalent for part-time workers) is −147 SEK, approximately 0.6 percent of an average monthly salary. The estimated reduction in weekly hours is precisely estimated, with a 95 percent CI from −0.80 to −1.77, whereas the monthly wage reduction is only marginally statistically distinguishable from zero (95 percent CI −312.6 to 17.9). Panels B and C of Figure 2 show the effect is quite stable over time for both wages and hours.

The modest wage response suggests a limited role for the wage margin in accounting for the overall labor supply response. To investigate the relative importance of the wage and hours margins more formally, we decompose the change in wage earnings into an hours and a wage component. Let $w_{i,t}$ denote the hourly wage and let $h_{i,t}$ denote annual hours worked by individual i at time t . The difference in wage earnings between time t and the year before the lottery can be written as

$$(3) \quad w_{i,t}h_{i,t} - w_{i,-1}h_{i,-1} = w_{i,t}h_{i,-1} + w_{i,-1}h_{i,t} + (w_{i,t} - w_{i,-1})(h_{i,t} - h_{i,-1}).$$

We estimate the contribution of changes on the wage and hours margin by using each of the three components on the right-hand side in (3) as dependent variables in regression (2) while controlling for $w_{i,-1}h_{i,-1}$. The five-year estimate indicates the reduction in hours worked accounts for 81 percent of the fall in wage earnings, whereas 18 percent is due to the negative effect of lottery wealth on wages, and only 1 percent to the interaction between hours and wages. Figure 2, panel D, shows the hours component dominates the wage component at all time horizons.

In online Appendix Section 5, we report on a number of robustness checks using contracted hours and alternative ways to impute earnings-based hours and wages. While these analyses indicate the hours component plays a relatively smaller role for the long-term earnings response, the hours component still dominates the wage effect at all time horizons.

Finally, we examine whether wealth affects employer, workplace, occupation, industry, or location of work. These variables are available for all employees, except occupation which is only available for a subset of employees from 1996 and onward. We find no evidence that wealth affects any of these variables in our analysis of five-year outcomes, nor in flexible analyses of the response at $t = 0, 1, \dots, 10$ (see online Appendix Figure A3). Because a plausible mechanism behind wage adjustments is that workers switch occupations, industries, or regions of work, the fact that we find no evidence of such switches is consistent with the hypothesis that changes in hours worked are likely to account for the bulk of the intensive margin response.

²²Imputing contracted hours from adjacent years does not mitigate the selection problem. To see this point, consider a worker who is covered by the survey in year t but quits the labor force in year $t + 1$. Imputing contracted hours in year $t + 1$ from year t implies we overstate the number of hours worked in $t + 1$.

In summary, we conclude that both extensive- and intensive-margin adjustments account for the responses we observe, and that wages contribute modestly to the adjustment on the intensive margin.

C. Heterogeneous and Nonlinear Effects

We conduct a number of analyses to examine whether the effects of wealth on our primary earnings measure are heterogeneous by lottery, sex, age at the time of win, education, pre-lottery earnings, and self-employment status. Figure 3 reports the labor supply trajectories for the different subsamples (except self-employment).²³

Figure 3, panel A, shows the effect is similar across lotteries, and we cannot reject the null hypothesis that the five-year estimates for the four lotteries are equal. Of particular interest is the comparison between Triss-Lumpsum and Triss-Monthly, because the underlying populations are the same, but the mode of payment differs. If winners have a significant bias to the present (O'Donoghue and Rabin 1999) and Triss-Monthly winners are unable to borrow against their future income stream, we would expect bigger immediate responses from lump-sum prizes. Yet the response patterns for the two Triss lotteries are quite similar, suggesting winners' behavior is consistent with a forward-looking dynamic labor supply model (which we estimate in the following section).

Standard life-cycle models predict stronger wealth effects for older workers because they have fewer years to spend the lottery prize. We test for heterogeneous effects by dividing the sample into three age ranges: 21–34, 35–54, and 55–64. As Figure 3, panel B, shows, the effects are similar by age in the years following the win. We fail to reject the null hypothesis that the five-year coefficients from the three subsamples are equal. Yet because the oldest age group has lower pre-win earnings, their response is larger relative to baseline (–8.9 percent of average pretax earnings for each 1 million SEK) compared to winners aged 21–34 (–5.9 percent) and 35–54 (–4.4 percent). Over longer time horizons, the effect tends to be weaker in the subsample of individuals in the 55–64 bracket, but this result is due to many of these individuals reaching retirement age, which mechanically attenuates the effect.

A common finding in the literature is that labor supply elasticities are larger for women than men (Keane 2011), though some recent work finds evidence of a decrease in labor supply elasticities for married women between the 1980s and 1990s (Blau and Kahn 2007). Our five-year estimates suggest that, if anything, women's labor supply responses to wealth shocks are weaker than those of men. The difference between the five-year estimates is not statistically significant ($p = 0.11$), and even if the coefficients are scaled relative to mean annual earnings (which are 31 percent lower for women), the coefficient estimates are in the opposite direction of what prior work typically has found. Yet the flexible coefficient estimate for $t = 0, 1, \dots, 10$, plotted in Figure 3, panel C, suggests the difference becomes smaller with time from the lottery. We do not infer from these results that women's labor supply is less responsive to wealth shocks than men's, but the 95 percent

²³The corresponding five-year estimates are reported in online Appendix Table A3.

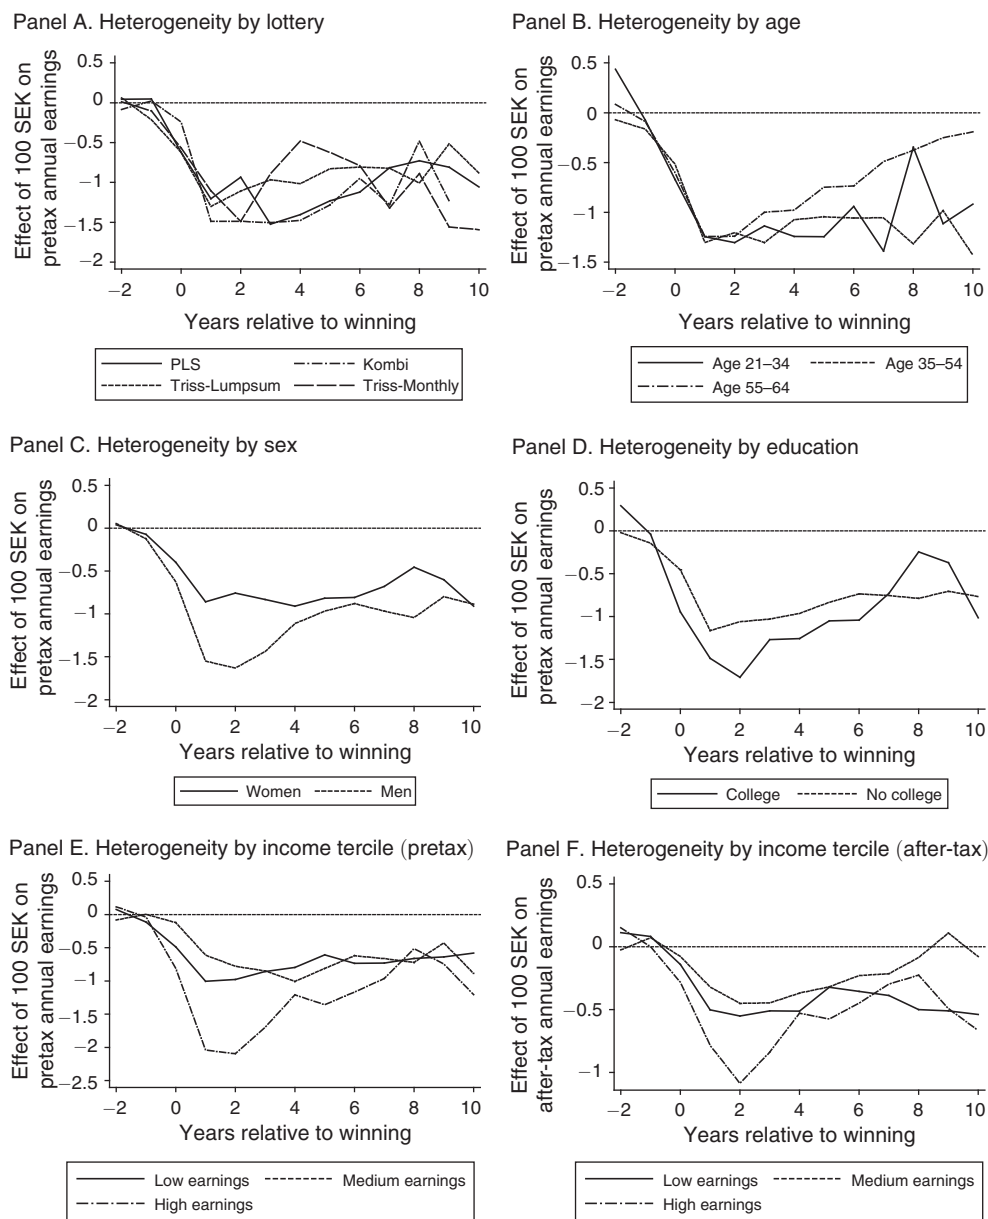


FIGURE 3. HETEROGENEOUS EFFECTS OF WEALTH ON EARNINGS

Notes: This figure reports estimates obtained from equation (2) estimated in different subsamples. The dependent variable is pretax labor earnings in panels A–E and after-tax labor income in panel F. A coefficient of 1.00 corresponds to an increase in annual earnings of 1 SEK for each 100 SEK won. Each year corresponds to a separate regression. The estimate for year 10 in panel A is excluded for Kombi winners because very few observations are available.

confidence intervals for the five-year estimates allow us to rule out that the female labor supply response exceeds the male response by more than 9 percent.

Panels E and F of Figure 3 show both the initial pretax and after-tax response is stronger for winners in the highest tercile of pre-lottery earnings, though we can only

marginally reject that the five-year estimates differ across income groups for pretax earnings ($p = 0.079$).

Earlier research has suggested the self-employed have greater flexibility in choosing their hours (Gurley-Calvez, Biehl, and Harper 2009; Hurst and Pugsley 2011). Yet the five-year estimates for self-employed (-1.130) and wage earners (-1.059) are very similar. We also find no evidence of heterogeneous effects depending on college completion.

Some theories predict wealth should have nonlinear effects on labor supply if workers who wish to reduce their labor supply face fixed adjustment costs (as in Chetty et al. 2011). In this case, the marginal effects of modest wealth shocks will be smaller than those of more substantial wealth shocks. We therefore estimate both a quadratic model and a spline model with a knot at 1 million SEK. The point estimates suggest the marginal effect of lottery wealth is smaller for larger prizes, but the difference is not statistically significant. Moreover, the estimated effect is about 10 percent to 30 percent larger when we exclude very large (≥ 5 million SEK), large (≥ 2 million), or moderate (≥ 1 million SEK) prizes.²⁴

IV. Dynamic Labor Supply Model

In this section, we estimate a simple dynamic life-cycle labor supply model using a simulated minimum-distance procedure. We use the model to recover a model-based estimate of the long-run, lifetime effect of a lottery prize on after-tax labor earnings, and to obtain estimates of key labor supply elasticities.

A. Model Setup

The model is a discrete-time, dynamic labor supply model with perfect foresight, no uncertainty, and no liquidity constraints. The agent lives for T periods ($t = 0, 1, \dots, T - 1$) and receives unearned income α_t in period t . Each period, the agent chooses consumption c_t , annual work hours h_t , and next period's assets (A_{t+1}). Annual earnings (y_t) are the product of the after-tax wage w_t and annual hours. Assets earn interest rate r between periods. Individuals in the model will choose to save for retirement, which must occur at $t = R^*$ or earlier; at this time, individuals can no longer choose $h_t > 0$.

Individuals make consumption, labor supply, and savings/borrowing decisions to maximize lifetime present discounted utility (using a discount rate δ), according to

$$(4) \quad U = \sum_{t=0}^{T-1} \frac{1}{(1+\delta)^t} \left(\beta \log(c_t - \gamma_c) + (1 - \beta) \log(\gamma_h - h_t) \right),$$

$$A_{t+1} = (1 + r)(A_t - c_t + w_t h_t + \alpha_t),$$

$$A_T \geq 0,$$

$$h_t = 0 \quad \text{for all } t \geq R^*.$$

²⁴Detailed results for the analysis of nonlinear effects are reported in online Appendix Table A4.

Following Bover (1989) and Imbens, Rubin, and Sacerdote (2001), we use a Stone-Geary utility function. The parameter β is the relative weight on consumption in utility, γ_c is the subsistence term for consumption, and γ_h is the maximum annual hours of work available. A lump-sum lottery prize is represented as a one-time shock to A_t . The empirical results provide individual-level estimates of $\partial y_{t+s} / \partial A_t$ for each time period following the lottery win.

We use the model to recover estimates of the lifetime MPE out of unearned income as well as uncompensated (Marshallian), compensated (Hicksian), and intertemporal (Frisch) labor supply elasticities. Before describing the simulation strategy, we discuss the role of three important model assumptions.

No Barriers to Saving and Borrowing.—We assume agents can save and borrow at interest rate r . An implication of this assumption is that two prizes with identical present discounted values should have the same dynamic effects on labor earnings. This model prediction is consistent with our reduced-form analysis, which finds similar results for Triss-Lumpsum and Triss-Monthly prizes.

Stone-Geary Functional Form.—Stone-Geary preferences simplify the simulation because the per-period problem can be solved in closed form. Additionally, in a static model, this functional form delivers an income effect that does not vary with the wage, which is consistent with our reduced-form finding that the after-tax earnings response is quite similar in different income groups.

Binding Retirement Age.—The Swedish retirement system admits flexibility in the timing of retirement, but as we discuss further in Section 6 in the online Appendix, a binding retirement age at 65 is a reasonable simplifying assumption. Clear “bunching” of retirement ages occurs at around age 65, with some retirement before age 65, but very little retirement after age 65. The model predicts no change in retirement age in response to small to modest lottery prizes because individuals prefer to smooth leisure and consumption over the life cycle. In this case, the binding retirement age is still binding after the lottery win. However, because the marginal utility of leisure is strictly positive even at zero hours of work, winners (of any age) may retire immediately and abstain from work completely if the amount won is sufficiently large.

B. Model Simulation

We simulate the model to match the main individual-level after-tax results. The years of life remaining depend on the age of the winner when the prize is awarded. When simulating the model, we match the empirical distribution of the age of winners in the data. Individuals retire at age 65 and die at age 80, so a 25-year-old winner would face $T = 55$ and $R^* = 40$. We choose $r = 0.02$ to match the average real risk-free rate in Sweden during the time period the data span. We assume the subsistence consumption term is $\gamma_c = 20,000$ SEK, about 12 percent of a median annual disposable income, and we assume the maximum annual hours of work available are $\gamma_h = 1,880$, which is the annual hours for a full-time worker in Sweden (working 40 hours per week with 5 weeks of mandated vacation). We set the wage in each period to be equal to the average after-tax labor income divided by average hours

worked in our data. Unearned income a_t is set to 0 for all $t < R^*$ and to 70 percent of average annual after-tax labor income for $t \geq R^*$.²⁵

We estimate via simulation the two remaining parameters, the discount rate (δ) and the relative weight on consumption in utility (β). For a given value of r , the time path of the labor earnings response following the lottery helps pin down δ . The lifetime earnings reduction to winning the lottery is primarily determined by the value of β , because this parameter governs the strength of the income effect.²⁶

We estimate the two parameters using a standard simulated minimum-distance procedure. For each set of parameters, we simulate the model and compute the effect of winning the lottery, i.e., $(\partial(y_t)/\partial(A_t), \dots, \partial(y_{t+10})/\partial(A_t))$.²⁷ We calculate these statistics for each simulated individual and then average across individuals, weighting individuals so that the age distribution in the simulated sample matches the lottery sample. See online Appendix Section 7 for further details about the model simulation, including the minimum-distance criterion we use and how we estimate standard errors for the parameter estimates.

C. Simulation Results and Implied Labor Supply Elasticities

Table 5 summarizes the simulation results. The χ^2 goodness-of-fit test statistic is not large ($\chi^2(8) = 3.433$, $p = 0.096$), suggesting the model provides a reasonably good fit to the reduced-form results. Additionally, the implied average annual hours are close to the average annual hours in the lottery sample (1,656 hours versus 1,633 hours). The estimate of β is 0.867 (SE = 0.046), suggesting (holding the marginal utility of wealth constant) roughly 13 percent of unearned income is spent reducing after-tax labor income, with the rest spent increasing consumption. The estimate of δ is 0.015 (SE = 0.037), which is close to the assumed interest rate of $r = 0.020$. This finding is consistent with fairly similar earnings responses over time, with the attenuation primarily driven by the mechanical effect of workers gradually reaching the binding retirement age.

Figure 4 compares the simulated model to the reduced-form effects of lottery wealth on after-tax income. Consistent with the relatively low χ^2 test statistic, the model-based estimates track the empirical estimates fairly closely. Panel B of Table 5 compares simulated results with empirical results that were not directly targeted in estimation, focusing on differences in the after-tax response by age, size

²⁵We make the simplifying assumption that pension income does not respond to labor earnings prior to retirement.

²⁶This discussion of identification is meant to convey intuition, but the actual identification of δ and β is more subtle. First, the lifetime earnings reduction is affected both by δ and β . Holding constant r and β , higher values of δ will increase the lifetime earnings reduction. Second, the binding retirement age will cause earnings reductions to decline over time mechanically as winners reach the binding retirement age. Thus, both a binding retirement age and $\delta > r$ will work toward producing reductions in annual earnings in the short run that are larger than the long-run reductions. Therefore, the full structure of the model is needed to separate the mechanical effect of retirement from the effect of the magnitude of δ relative to r .

²⁷To compute this effect, we first solve a full life-cycle model assuming no lottery win and perfect foresight. Then, this solved model (solved over a large grid of possible asset choices) also contains the implied solution of how households would re-optimize at whatever age the person wins the lottery. In other words, the lottery is treated as an unexpected shock to assets at some time t , and so up until the time of the lottery prize, the individual follows the “no lottery” optimal path of asset accumulation, and then assets jump from A_t to the $A_t + L$ and then the dynamic programming solution gives a new (re-optimized) path of labor earnings, consumption, and savings following the lottery win for the remaining time periods.

TABLE 5—SIMULATION-BASED ESTIMATES OF MODEL PARAMETERS

Panel A. Parameter estimates			Panel B. Model fit		
	Estimate	SE		Reduced form	Model prediction
Consumption weight (β)	0.867	(0.046)	Baseline	−0.58	−0.55
Discount rate (δ)	0.015	(0.037)	High wage	−0.58	−0.55
	$\chi^2(8)$	p -value	Low wage	−0.54	−0.56
Goodness of fit	3.433	[0.096]	100,000 SEK prize	−0.54	−0.55
			2 million SEK prize	−0.55	−0.55
	Data	Model	Age 21–34	−0.73	−0.38
Average annual hours	1,633.0	1,655.8	Age 35–54	−0.68	−0.53
			Age 55–64	−0.38	−0.58
Panel C. Implied wealth effect by age			Panel D. Implied labor supply elasticities		
Assumed age-at-win	Lifetime MPE		Marshallian Elasticity (e_M)	0.009	
20	−0.168		Hicksian Elasticity (e_H)	0.095	
30	−0.144		Frisch Elasticity (e_F)	0.143	
40	−0.118				
50	−0.086				
60	−0.036				

Notes: Panel A reports results of estimating the dynamic model via simulated method of moments. The goodness-of-fit test uses the minimized value of weighted minimum distance procedure, based on ten moments and two parameters. Panel B compares the model-generated predictions to our causal estimates; in each case, the comparison is to the five-year estimate for after-tax labor income. Panel C reports the lifetime wealth effect at different ages at the time of win. Panel D reports key labor supply elasticities implied by the model-generated parameters for individuals who play the lottery at age 50. In these analyses, we assume individuals retire at 65 and die at age 80. Panel C reports the effect of a lottery prize on total labor earnings (i.e., sum of $\partial y / \partial L$ across all remaining working years, as implied by model), whereas panel D reports the implied effect of a permanent increase of wages on total hours worked (summed up across all remaining working years), the implied Hicksian elasticity (calculated from the Slutsky equation), and the Frisch elasticity.




FIGURE 4. COMPARING MODEL-BASED ESTIMATES TO EMPIRICAL RESULTS

Notes: This figure compares the model-based estimates using the best-fit parameters reported in Table 5 to the estimates obtained from equation (2) estimated in the pooled lottery sample with after-tax labor income as the dependent variable. Each year corresponds to a separate regression. The graph shows results for $t = -1$ for illustrative purpose only. We control for earnings at $t = -1$, so the empirical estimate is exactly zero. In the simulation, the prize is assumed to be awarded at end of year 0, so $\partial y / \partial L$ for both $t = -1$ and $t = 0$ are zero by assumption.

of prize amount, and pre-win earnings of the winner. Our simulation results are broadly in line with the empirical results, which show fairly limited variation across pre-win earnings and the size of the prize. For age, the results are mixed. The empirical results indicate smaller estimates for older winners, although the differences by age are not statistically significant. By contrast, the results by age in the simulated model indicate the opposite pattern. This finding suggests other important factors may be present, such as human-capital accumulation, that are outside the model, but are important for understanding differences in wealth effects by age.²⁸

Estimating the Lifetime MPE.—Using the estimates of the model, we can compute the lifetime MPE (after-tax) income out of unearned income, where the calculation extrapolates beyond the first ten years following the lottery win to the entire remaining years of life. The model estimates imply lifetime MPEs that vary with age at the time of win, from -0.17 for 20-year-old winners to -0.04 for winners aged 60 (see panel C of Table 5). For younger winners, the model estimates imply most of the lifetime-earnings reduction occurs after the first ten years, implying the cumulative ten-year effects significantly understate lifetime wealth effects. Estimates of the MPE previously reported in the literature vary substantially, but the average lifetime MPE in our data (-0.08) is lower than the median (-0.15) among the 30 different estimates reported by Pencavel (1986).²⁹ Incidentally, our average MPE is closer to the MPE of -0.11 reported by Imbens, Rubin, and Sacerdote (2001) when they exclude non-winners and winners of extremely large prizes from their data.³⁰

Recovering Key Labor Supply Elasticities.—Using the full structure of the model, we can also recover key labor supply elasticities that feature prominently in previous research. In panel D of Table 5, we report the uncompensated (Marshallian) elasticity, the compensated (Hicksian) elasticity, and the intertemporal (Frisch) elasticity. The simulated elasticities are computed for someone who wins at age 50. The uncompensated elasticity is very small in magnitude, which is a direct consequence of the Stone-Geary functional-form assumption. The Hicksian elasticity is estimated to be around 0.10, which is smaller than the average Hicksian elasticity estimate of 0.31 reported in the meta analysis in Keane (2011).

²⁸Despite the many simplifying assumptions, we note the model can also provide a reasonable fit for asset accumulation over the life cycle in a Swedish representative sample. Online Appendix Figure A4 shows the simulated asset path for a 25-year-old non-winner together with the median and mean net wealth by age in a Swedish representative sample in year 2000. The simulated model assumes lifespan ends at 80 and no bequest motive exists; either a bequest motive or uncertain lifespan would likely allow the model to better fit the wealth data after age 65.

²⁹Two recent studies that consider settings similar to ours find substantially larger MPEs than we do. Kimball and Shapiro (2008) estimate an MPE of -0.37 using survey responses about hypothetical lottery winnings, whereas Bengtsson (2012) estimates an MPE of about -0.30 among recipients of unconditional cash grants in South Africa.

³⁰The similarity in terms of average MPEs masks nontrivial differences in estimation and modeling. Plugging our five-year estimate for the after-tax response (column 10 of Table 3) into the model in Imbens, Rubin, and Sacerdote (2001) gives an MPE of -0.05 . The reason for the lower MPE is that they assume $\delta = r = 0.10$, whereas we assume $r = 0.02$ and estimate δ to be 0.015. A high interest rate implies lump-sum prizes are “large” relative to yearly installments (the setting studied by Imbens, Rubin, and Sacerdote 2001), attenuating the MPE based on our estimates. The same exercise with $\delta = r = 0.02$ gives an MPE based on our estimates of -0.13 compared to -0.14 based on the estimates in Imbens, Rubin, and Sacerdote (2001). The reason for the higher MPE compared to our calibration is the high implicit retirement age in Imbens, Rubin, and Sacerdote (2001). Because they assume winners continue working for 30 years, the implicit average retirement age would be 78 in our sample and 80 in theirs.

The Frisch elasticity is estimated to be close to 0.14, which is smaller than the range of estimates (0.27–0.53) used by the CBO (Reichling and Whalen 2012). Although these specific elasticities are recovered from the reduced-form income-effect estimates and the functional-form assumptions of the dynamic model, the specific Stone-Geary functional form does not entirely drive the estimated elasticities. In a wide range of time-separable utility models, the Frisch elasticity and the Hicksian elasticity are related by the intertemporal elasticity of substitution (IES), the estimated income effect, and the ratio of wealth to income (Ziliak and Kniesner 1999; Browning 2005). Therefore, modest estimates of the income effect necessarily constrain the Frisch elasticity to be similar in magnitude to the Hicksian elasticity, as long as the IES and the Marshallian elasticity are not very large in magnitude.³¹ We illustrate this point through a series of sensitivity analyses (reported in online Appendix Section 7) that report broadly similar elasticities under different assumptions on the interest rate, consumption floor, the IES, and the Marshallian elasticity.

V. Household-Level Analyses

Two questions guide our household-level analyses. First, if winners' spouses also adjust their labor supply following a wealth shock, individual-level estimates will understate the overall labor supply response, implying elasticities inferred under the assumption that the winner's response fully captures the labor supply effects of the wealth shock are potentially misleading. Because the register data contain the spouses of winners, we can test for and quantify the size of the difference between the household- and individual-level responses.

Second, we use our data to test the unitary model of the household, in which two spouses are modeled as a single decision-making unit (Becker 1973; Becker 1976). These models make the strong prediction that the identity of a spouse who experiences a random wealth shock should not influence the labor supply responses of each of the two spouses (see Lundberg, Pollak, and Wales 1997 and Attanasio and Lechene 2002 for similar empirical tests).

We conduct our household-level analyses by augmenting the sample of married individuals with their spouses. The key results are summarized in Table 6. Beginning with our first question, columns 1–3 of panel A shows the five-year estimates for pretax labor earnings of married winners, spouses, and married households (defined as the sum of the winner's and spouse's labor supply response). Figure 5 shows the corresponding dynamic effects.

We find that married winners reduce their pretax annual labor earnings by 0.97 SEK per 100 SEK won, compared to 0.41 SEK for their spouses. The total

³¹If lifetime utility is additively separable, and there is perfect foresight, no uncertainty, and perfect capital markets, the relation between the Frisch (e_F) and the Hicksian (e_H) elasticity is $e_F = e_H + \rho(\partial(wh)/\partial A)^2(A/w_h)$, where ρ is the IES, $\partial(wh)/\partial A$ is the income effect, and A/w_h is the ratio of wealth to income (see Ziliak and Kniesner 1999; and Browning 2005). In the calculations in panel D of Table 5, e_H is roughly 0.1, ρ is roughly 1 given Stone-Geary utility, the income effect is roughly 0.09, and the ratio of A/w_h is approximately 6.7. This implies an estimate of e_F of 0.15, which is very close to the value calculated directly from model simulation. Assuming a small Marshallian elasticity, e_H and the income effect will be similar in magnitude from the Slutsky equation. A large Frisch elasticity consequently requires a large value of IES. A doubling of IES to 2.0 and an increase in magnitude of Marshallian elasticity to 0.2 would still give a value of Frisch elasticity below 0.4.

TABLE 6—EFFECT OF WEALTH ON HOUSEHOLD EARNINGS

	Married lottery players			Unmarried lottery players	Total household effect
	Winner (1)	Spouse (2)	Household (3)	players (4)	effect (5)
<i>Panel A. Individual and household labor supply responses</i>					
Effect (100 SEK)	−0.965	−0.408	−1.373	−1.289	−1.306
SE	(0.201)	(0.208)	(0.299)	(0.228)	(0.194)
<i>p</i>	[<0.001]	[0.049]	[<0.001]	[<0.001]	[<0.001]
Mean (SEK)	195,848	186,017	381,865	195,034	301,330
Effect/mean	−4.93	−2.19	−3.60	−6.61	−4.33
Observations	142,102	142,102	142,102	99,287	244,826
PLS and Kombi					
				Both spouses age ≤ 64	
	Married lottery players (6)	(7)	(8)	Spouse ever played lottery (9)	Spouse played in draw (10)
<i>Panel B. Difference between winner and spousal responses</i>					
Effect (100 SEK)	−0.557	−0.964	−0.812	−1.241	−1.338
SE	(0.278)	(0.397)	(0.422)	(0.588)	(0.624)
<i>p</i>	[0.045]	[0.015]	[0.054]	[0.035]	[0.032]
Number of PLS tickets winner/spouse		173.59/48.6	168.9/47.4	162.1/73.7	164.1/113.6
Number of Kombi tickets winner/spouse		1.34/0.15	1.34/0.14	1.33/0.97	1.34/1.20
Observations	142,102	140,398	127,344	75,893	49,352

Notes: This table reports five-year estimates obtained by estimating equation (2) on winners, winners' spouses, and at the household level for different subsamples with pretax labor earnings as the dependent variable. Panel A shows the effect on winners, spouses, and households, whereas panel B shows the estimated differential response for winners relative to spouses. Except in columns 4 and 5, the estimation sample is restricted to married couples. Results in columns 7–10 are restricted to the PLS and Kombi lotteries. In column 8, we further restrict the sample to couples in which both the winner and the spouse were between 21 and 64 at the time of the win. The prize amount is scaled so that a coefficient of 1.00 implies a 1 SEK increase in earnings per 100 SEK won. The estimates include baseline controls for the winner's spouse (when applicable). The number of lottery tickets refers to the average number of tickets held by winners in each lottery (and Kombi controls) and their respective spouses. Effect/mean: effect of 1 million SEK relative to the mean of the dependent variable (in percent). Standard errors are clustered at the level of the player.

household-level response of -1.373 is thus substantially stronger than the individual-level response of married individuals. Column 4 of Table 6 shows earnings of unmarried winners fall 1.29 SEK per 100 SEK won, more than for married winners, but less than the household-level response for married couples. Finally, column 5 shows the effect on household labor supply for the full sample. Including the response of nonwinning spouses increases the labor supply response from -1.066 (column 1 of Table 3) to -1.306 . Including the spousal response thus increases the estimated labor supply response by 23 percent.

Turning to the second question, panel B of Table 6 shows the difference between the labor supply responses of winners and spouses. Negative estimates imply the winner reacts more strongly than the spouse. Column 6 shows the difference in the full sample (i.e., between the labor supply response of winners and spouses in columns 1 and 2). Married winners reduce their labor supply by 0.56 SEK more than their spouses for every 100 SEK won ($p = 0.045$), a finding seemingly at odds with income pooling.




FIGURE 5. EFFECT OF WEALTH ON EARNINGS OF MARRIED WINNERS AND SPOUSES

Notes: This figure reports estimates obtained from estimating equation (2) separately for married winners, their spouses, and married households. The dependent variable is pretax labor earnings. Each year corresponds to a separate regression.

To more carefully assess the unitary model, we exclude the Triss lottery from columns 7–10, for two reasons. First, married couples may sometimes buy Triss lottery tickets together, implying ownership of the winning ticket within the couple is unclear.³² By contrast, both the winning account in PLS and lottery ticket subscription in Kombi pertain to a specific individual. Using data from the Wealth Registry, we find married winners in Kombi and PLS retain a larger share of households' observable lottery wealth (78 percent and 85 percent) than married Triss winners (72 percent), suggesting within-couple ownership is indeed more clearly defined in the former two lotteries.³³ Second, nonwinning spouses may differ systematically from winning spouses in ways that correlate with how they respond to wealth shocks. In Triss, this concern is difficult to put to a stringent test, because we do not have information about the population of lottery players who selected into the lottery, only players who appear on the TV show. In PLS and Kombi, we have information about the universe of players and the number of tickets owned. This information allows us to test if the differential response observed between winners and their spouses persists in households where both spouses participated in the lottery.

Column 7 of Table 6 shows restricting attention to the PLS and Kombi samples increases the spousal difference to -0.964 ($p = 0.015$), in line with the relatively

³²The Triss data contain information about shared ownership of lottery tickets, but the data rarely indicate shared ownership between married spouses, probably because "contracts" regarding ownership are less explicit between spouses, and because wealth is split equally in the event of a divorce. Consequently, in some cases, married couples are likely to have bought a winning ticket together, but only one of the spouses appears on the show.

³³Online Appendix Figure A5 and Table A5 show the complete results for how lottery wealth is allocated between spouses. Because the Swedish Wealth Registry only existed in 1999–2007, we observe wealth for very few winners in PLS and therefore use capital income as a proxy for wealth in this case. We exclude Triss-Monthly winners because inferring how the prize money is allocated within couples when it is paid out over a long time is difficult.

larger share of the wealth shock that pertains to the winner in PLS and Kombi. Column 8 shows the difference decreases somewhat when we further restrict the sample to couples in which both spouses (and not just the winner) were below the age of 64 at the time of win (-0.812). We impose this restriction because retired spouses may be constrained in their labor supply choices.

Next, we attempt to reduce any biases due to possible nonrandomness in which spouse wins the lottery. In column 9, we restrict the sample to couples in which the nonwinning spouse participated in the winning draw or pre-win draws in the same lottery. In column 10, we go further and restrict the sample to couples in which both spouses participated in the winning draw. Imposing these sample restrictions reduces the difference between winners and spouses both in terms of number of lottery tickets held (see Table 6) and demographic characteristics (online Appendix Table A6). It is therefore reassuring that imposing these restrictions strengthens the differential response between winners and spouses.³⁴ It is also reassuring that the winner's share of the households' total labor supply response in columns 7 to 10 (between 79 and 89 percent) corresponds well with the share of lottery wealth allocated to the winner in PLS and Kombi.

In additional analyses, we find no clear evidence that the effect of lottery wealth on winner and spousal earnings depends on the winner's sex or whether the primary or secondary earner wins the lottery.³⁵ Because of the smaller sample sizes, however, these estimates are considerably less precise. Another concern with our household-level results is that lottery wealth might affect household composition. In online Appendix Section 9, we find a small positive, but statistically insignificant, effect of lottery wealth on divorce risk. Our results do not change appreciably when the sample is restricted to couples that remain married.

We conclude from our household analyses that estimates of wealth effects on married individuals' labor supply underestimate the overall household labor supply response. Across a suite of analyses, we also consistently find the winning spouse responds more strongly than the nonwinning spouse, with the strongest evidence in the specifications that most closely approximate the ideal experiment.

The household-level results suggest the identity of the winner determines who in a married couple reduces labor supply the most, which is inconsistent with the unitary model. A prominent alternative class of non-unitary household models emphasizes the role of bargaining under the threat of divorce (Manser and Brown 1980; McElroy and Horney 1981). Reconciling our results with this class of models, however, is difficult. According to Swedish marriage law, the default rule in the event of divorce is that all assets are divided equally between spouses, unless the couple has a prenuptial agreement. Prenuptial agreements are uncommon and lottery winnings will therefore, in most cases, affect the outside option of the winner and spouse symmetrically.³⁶

³⁴Consistent with the weaker bargaining position of cohabitants, we find the estimated difference between winners' and spouses' responses becomes stronger if we define cohabiting couples with children as "married" (identifying cohabiting couples without children is not possible in the Swedish data).

³⁵We report these analyses in online Appendix Table A7. When including the Triss sample, we obtain suggestive evidence that the differential response is stronger when the husband or the primary earner wins the lottery.

³⁶Online Appendix Section 8 provides additional details about Swedish marital law and documents that prenuptial agreements are uncommon.

Instead, our results appear to be more consistent with the “separate spheres” household model of Lundberg and Pollak (1993) that relies on bargaining with threat points internal to the marriage. As long as a couple remains married, the winner owns and controls the prize money unless he or she decides to transfer part of the prize to the nonwinning spouse, or deposit the money in a joint account. Lottery wealth can therefore improve the bargaining power of the winner by making the winner better off in the within-marriage noncooperative equilibrium that defines the threat points in the separate-spheres model. Our results are consistent with related work that finds empirical support of the separate-spheres model over divorce-threat models (Lundberg, Pollak, and Wales 1997; Lundberg, Startz, and Stillman 2003).³⁷

VI. Conclusion

We have shown that an exogenous wealth shock results in an immediate and permanent reduction in earnings. The magnitude of the response is modest; pretax earnings decrease by about 1 percent of the wealth shock in each of the first ten years following the win. The response is about 40 percent smaller when we instead consider after-tax income, and about 40 percent larger when we measure labor supply in terms of production value (earnings including employer-paid social security contributions). The earnings response is mainly due to a reduction in wage earnings due to fewer hours worked. Although adjustment takes place on both the extensive and intensive margin, our estimates suggest the latter is more important.

A surprising finding is the limited heterogeneity across many interesting demographic subgroups. Imbens, Rubin, and Sacerdote (2001) similarly find no significant differences in the responses of men and women, and note this finding is at odds with a large literature that finds women are systematically more responsive to price and wealth changes. We extrapolate our main results using a dynamic life-cycle model in order to calculate lifetime wealth effects, which we find vary with age and are larger for younger workers. Because the estimated wealth effects are modest, our calibrated model implies labor supply elasticities that are in the lower range of previously reported estimates. In our household-level analyses, we find both winners and spouses reduce their labor supply, but the reduction is stronger for winners. This finding provides unusually strong evidence against the testable prediction of unitary models of household labor supply that exogenous unearned income is pooled within the household.

REFERENCES

- Andersen, Steffen, and Kasper Meisner Nielsen 2012. “Ability or Finances as Constraints on Entrepreneurship? Evidence from Survival Rates in a Natural Experiment.” *Review of Financial Studies* 25 (12): 3684–710.

³⁷ Both the “divorce threat” and “separate spheres” models are cooperative household models. Our results may also be consistent with specific noncooperative models (Lundberg and Pollak 1994; Doepke and Tertilt 2014), as well as limited commitment models that combine elements of cooperative and noncooperative models (Ligon, Thomas, and Worrall 2002; Voena 2015). Distinguishing between these models is difficult and subtle, which is why we focus on testing the income-pooling property of the unitary model and comparing two specific non-unitary cooperative models.

- Attanasio, Orazio, and Valérie Lechene. 2002. "Tests of Income Pooling in Household Decisions." *Review of Economic Dynamics* 5 (4): 720–48.
- Becker, Gary S. 1973. "A Theory of Marriage: Part I." *Journal of Political Economy* 81 (4): 813–46.
- Becker, Gary S. 1976. *The Economic Approach to Human Behavior*. Chicago: University of Chicago Press.
- Bengtsson, Niklas. 2012. "The Marginal Propensity to Earn and Consume out of Unearned Income: Evidence Using an Unusually Large Cash Grant Reform." *Scandinavian Journal of Economics* 114 (4): 1393–413.
- Blau, Francine D., and Lawrence M. Kahn. 2007. "Changes in the Labor Supply Behavior of Married Women: 1980–2000." *Journal of Labor Economics* 25 (3): 393–438.
- Blundell, Richard, and Thomas MaCurdy. 2000. "Labor Supply: A Review of Alternative Approaches." In *Handbook of Labor Economics*, Vol. 3A, edited by Orley Ashenfelter and David Card, 1559–695. Amsterdam: Elsevier.
- Bodkin, Ronald. 1959. "Windfall Income and Consumption." *American Economic Review* 49 (4): 602–14.
- Bover, Olympia. 1989. "Estimating Intertemporal Labour Supply Elasticities Using Structural Models." *Economic Journal* 99 (398): 1026–39.
- Browning, Martin. 2005. "A Working Paper from April 1985: Which Demand Elasticities Do We Know and Which Do We Need to Know for Policy Analysis?" *Research in Economics* 59 (4): 293–320.
- Cesarini, David, Erik Lindqvist, Matthew Notowidigdo, and Robert Östling. 2017. "The Effect of Wealth on Individual and Household Labor Supply: Evidence from Swedish Lotteries: Dataset." *American Economic Review*. <https://doi.org/10.1257/aer.20151589>.
- Cesarini, David, Erik Lindqvist, Robert Östling, and Björn Wallace. 2016. "Wealth, Health, and Child Development: Evidence from Administrative Data on Swedish Lottery Players." *Quarterly Journal of Economics* 131 (2): 687–738.
- Chetty, Raj, John N. Friedman, Tore Olsen, and Luigi Pistaferri. 2011. "Adjustment Costs, Firm Responses, and Micro vs. Macro Labor Supply Elasticities: Evidence from Danish Tax Records." *Quarterly Journal of Economics* 126 (2): 749–804.
- Congressional Budget Office (CBO). 1996. *Labor Supply and Taxes*. Washington, DC: Congressional Budget Office.
- Doepke, Matthias, and Michèle Tertilt. 2014. "Does Female Empowerment Promote Economic Development?" National Bureau of Economic Research Working Paper 19888.
- Donni, Olivier, and Pierre-André Chiappori. 2011. "Non-unitary Models of Household Behavior: A Survey of the Literature." In *Household Economic Behaviors*, edited by Alberto Molina, 1–40. Berlin: Springer.
- Furåker, Bengt, and Anna Hedenus. 2009. "Gambling Windfall Decisions: Lottery Winners and Employment Behavior." *UNLV Gaming Research and Review Journal* 13 (2): 1–15.
- Gurley-Calvez, Tami, Amelia Biehl, and Katherine Harper. 2009. "Time-Use Patterns and Women Entrepreneurs." *American Economic Review* 99 (2): 139–44.
- Hankins, Scott, and Mark Hoekstra. 2011. "Lucky in Life, Unlucky in Love? The Effect of Random Income Shocks on Marriage and Divorce." *Journal of Human Resources* 46 (2): 403–26.
- Hankins, Scott, Mark Hoekstra, and Paige Marta Skiba. 2011. "The Ticket to Easy Street? The Financial Consequences of Winning the Lottery." *Review of Economics and Statistics* 93 (3): 961–69.
- Hirano, Keisuke, and Guido Imbens. 2004. "The Propensity Score with Continuous Treatments." In *Applied Bayesian Modelling and Causal Inference from Missing Data Perspectives*, edited by Andrew Gelman and Xiao-Li Meng, 73–84. New York: Wiley.
- Holtz-Eakin, Douglas, David Joulfaian, and Harvey Rosen. 1993. "The Carnegie Conjecture: Some Empirical Evidence." *Quarterly Journal of Economics* 108 (2): 413–35.
- Holtz-Eakin, Douglas, David Joulfaian, and Harvey Rosen. 1994. "Sticking It Out: Entrepreneurial Survival and Liquidity Constraints." *Journal of Political Economy* 102 (1): 53–75.
- Hurst, Erik, and Benjamin Wild Pugsley. 2011. "What Do Small Businesses Do?" National Bureau of Economic Research Working Paper 17041.
- Imbens, Guido W., Donald B. Rubin, and Bruce I. Sacerdote. 2001. "Estimating the Effect of Unearned Income on Labor Earnings, Savings, and Consumption: Evidence from a Survey of Lottery Players." *American Economic Review* 91 (4): 778–94.
- Joulfaian, David, and Mark O. Wilhelm. 1994. "Inheritance and Labor Supply." *Journal of Human Resources* 29 (4): 1205–34.
- Kaplan, H. Roy. 1985. "Lottery Winners and Work Commitment: A Behavioral Test of the American Work Ethic." *Journal of the Institute for Socioeconomic Research* 10 (2): 82–94.
- Keane, Michael P. 2011. "Labor Supply and Taxes: A Survey." *Journal of Economic Literature* 49 (4): 961–1075.

- Kearney, Melissa S., Peter Tufano, Erik Hurst, and Jonatahn Guryan. 2011. "Making Savings Fun: An Overview of Prize-Linked Savings." In *Financial Literacy: Implications for Retirement Security and the Financial Marketplace*, edited by Olivia Mitchell and Annamaria Lusardi, 218–40. Oxford: Oxford University Press.
- Kimball, Miles S., and Matthias D. Shapiro. 2008. "Labor Supply: Are the Income and Substitution Effects Both Large or Both Small?" National Bureau of Economic Research Working Paper 14208.
- Krueger, Alan B., and Jörn-Steffen Pischke. 1992. "The Effect of Social Security on Labor Supply: A Cohort Analysis of the Notch Generation." *Journal of Labor Economics* 10 (4): 412–37.
- Lagerwall, Björn. 2008. *Real Interest Rates in Sweden*. Stockholm: Sveriges Riksbank.
- Larsson, Bengt. 2011. "Becoming a Winner but Staying the Same: Identities and Consumption of Lottery Winners." *American Journal of Economics and Sociology* 70 (1): 187–209.
- Ligon, Ethan, Jonathan P. Thomas, and Tim Worrall. 2002. "Informal Insurance Arrangements with Limited Commitment: Theory and Evidence from Village Economies." *Review of Economic Studies* 69 (1): 209–44.
- Lindh, Thomas, and Henry Ohlsson. 1996. "Self-Employment and Windfall Gains: Evidence from the Swedish Lottery." *Economic Journal* 106 (439): 1515–26.
- Lundberg, Shelly, and Robert A. Pollak. 1993. "Separate Spheres Bargaining and the Marriage Market." *Journal of Political Economy* 101 (6): 988–1010.
- Lundberg, Shelly, and Robert A. Pollak. 1994. "Noncooperative Bargaining Models of Marriage." *American Economic Review* 84 (2): 132–37.
- Lundberg, Shelly, and Robert A. Pollak. 1996. "Bargaining and Distribution in Marriage." *Journal of Economic Perspectives* 10 (4): 139–58.
- Lundberg, Shelly, Robert A. Pollak, and Terrence J. Wales. 1997. "Do Husbands and Wives Pool Their Resources? Evidence from the United Kingdom Child Benefit." *Journal of Human Resources* 32 (3): 463–80.
- Lundberg, Shelly, Richard Startz, and Steven Stillman. 2003. "The Retirement-Consumption Puzzle: A Marital Bargaining Approach." *Journal of Public Economics* 87 (5–6): 1199–218.
- Manser, Marilyn, and Murray Brown. 1980. "Marriage and Household Decision-Making: A Bargaining Analysis." *International Economic Review* 21 (1): 31–44.
- McClelland, Robert, and Shannon Mok. 2012. "A Review of Recent Research on Labor Supply Elasticities." Congressional Budget Office Working Paper 2012-12.
- McElroy, Marjorie B., and Mary Jean Horney. 1981. "Nash-Bargained Household Decisions: Toward a Generalization of the Theory of Demand." *International Economic Review* 22 (2): 333–49.
- Mirrlees, James A. 1971. "An Exploration in the Theory of Optimum Income Taxation." *Review of Economic Studies* 38 (114): 175–208.
- O'Donoghue, Ted, and Matthew Rabin. 1999. "Doing It Now or Later." *American Economic Review* 89 (1): 103–24.
- Pencavel, John. 1986. "Labor Supply of Men: A Survey." In *Handbook of Labor Economics*, Vol. 1, edited by Orley Ashenfelter and Richard Layard, 3–102. Amsterdam: Elsevier.
- Picchio, Matteo, Sigrid Suetens, and Jan C. van Ours. 2017. "Labour Supply Effects of Winning a Lottery." *Economic Journal*. DOI: 10.1111/eoj.12465.
- Prescott, Edward C. 1986. "Theory Ahead of Business-Cycle Measurement." *Carnegie-Rochester Conference Series on Public Policy* 25: 11–44.
- Rebelo, Sergio. 2005. "Real Business Cycle Models: Past, Present and Future." *Scandinavian Journal of Economics* 107 (2): 217–38.
- Reichling, Felix, and Charles Whalen. 2012. "Review of Estimates of the Frisch Elasticity of Labor Supply." Congressional Budget Office Working Paper 2012-13.
- Saez, Emmanuel. 2001. "Using Elasticities to Derive Optimal Income Tax Rates." *Review of Economic Studies* 68 (1): 205–29.
- Saez, Emmanuel, Joel Slemrod, and Seth H. Giertz. 2012. "The Elasticity of Taxable Income with Respect to Marginal Tax Rates: A Critical Review." *Journal of Economic Literature* 50 (1): 3–50.
- Taylor, Mark P. 2001. "Self-Employment and Windfall Gains in Britain: Evidence from Panel Data." *Economica* 68 (272): 539–65.
- Voena, Alessandra. 2015. "Yours, Mine, and Ours: Do Divorce Laws Affect the Intertemporal Behavior of Married Couples?" *American Economic Review* 105 (8): 2295–332.
- Ziliak, James P., and Thomas J. Kniesner. 1999. "Estimating Life Cycle Labor Supply Tax Effects." *Journal of Political Economy* 107 (2): 326–59.